



Image and Vision Computing

Features



Instructor: Chen Yisong
HCI & Multimedia Lab, Peking University




实践出真知

纸上得来终觉浅
绝知此事要躬行


——陆游《冬夜读书示子聿》

古人家问无遗力，少壮工夫老始成。纸上得来终觉浅，绝知此事要躬行。



Today's Goals


- Features Overview
- Canny Edge Detector
- Harris Corner Detector
- Templates and Image Pyramid
- SIFT Features




Today's Questions

- What is a feature?
- What is an image filter?
- How can we find edges?
- How can we find corners?


- (How can we find cars in images?)



What is a Feature?



- Local, meaningful, detectable parts of the image




Features in Computer Vision

- What is a feature?
 - Location of sudden change
- Why use features?
 - Information content high
 - Invariant to change of view point, illumination
 - Reduces computational burden


Vanishing Points (无穷远点/灭点)



Vanishing Line (地平线)



Local versus global



Vanishing Line




Image 1








Image 2




Features in computer vision

- Compositing



This is your test image set


Features in computer vision




- Mosaic



Where Features Are Used



- Calibration(相机标定)
- Image Segmentation(图像分割)
- Correspondence in multiple images (对应匹配)
- Object detection, recognition(检测识别)




What Makes For *Good* Features?


- Invariance
 - View point (scale, orientation, translation)
 - Lighting condition
 - Object deformations
 - Partial occlusion
- Other Characteristics
 - Uniqueness
 - Sufficiently many
 - Tuned to the task

Invariant Local Features

Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters




SIFT = Scale Invariant Feature Transform




Advantages of local features

- **Locality:** features are local, so robust to occlusion and clutter (no prior segmentation)
- **Distinctiveness:** individual features can be matched to a large database of objects
- **Quantity:** many features can be generated for even small objects
- **Efficiency:** close to real-time performance
- **Extensibility:** can easily be extended to wide range of differing feature types, with each adding robustness




More motivation...

- Feature points are used also for:
 - Image alignment (图像配准/对齐)
 - 3D reconstruction(三维重构)
 - Motion tracking(运动跟踪)
 - Object recognition(目标识别)
 - Indexing and database retrieval(信息检索)
 - Robot vision(机器人视觉)
 - Others.....




Origin of Edges





- Edges are caused by a variety of factors



We also get: Boundaries of surfaces




Boundaries of depths






Boundaries of materials properties






Boundaries of lighting








Edge Detection




- Convert a 2D image into a set of curves
 - Extracts *salient features* of the scene
 - More compact than pixels





Edge Detection




- How can you tell that a pixel is on an edge?



Edge Types



Real Edges




Noisy and Discrete!

We want an **Edge Operator** that produces:

- Edge Magnitude
- Edge Orientation
- High Detection Rate and Good Localization




Edge Detection Continued




Boundary Detection – Edges


- Boundaries of objects
 - Usually different materials/orientations, intensity changes.



Edge is Where Change Occurs


- Change is measured by derivative in 1D
- Biggest change, derivative has maximum magnitude
- Or 2nd derivative is zero.






Noisy Step Edge

- Gradient is high everywhere.
- Must smooth before taking gradient.







So, 1D Edge Detection has steps:

1. Filter out noise: convolve with Gaussian
2. Take a derivative: convolve with $[-1 \ 0 \ 1]$
3. Find the peak. Two issues:
 - Should be a local maximum.
 - Should be sufficiently high.


What is the gradient?



What is the gradient?




What is the gradient?



Gradient direction is perpendicular to edge.

Gradient Magnitude measures edge strength.



Detecting Discontinuities

- Image derivatives


$$\frac{\partial f}{\partial x} = \lim_{\varepsilon \rightarrow 0} \left(\frac{f(x + \varepsilon) - f(x)}{\varepsilon} \right) \rightarrow \frac{\partial f}{\partial x} \approx \frac{f(x_{n+1}) - f(x)}{\Delta x}$$

- Convolve image with derivative filters

Backward difference [-1 1]

Forward difference [1 -1]

Central difference [-1 0 1]



Derivative in Two-Dimensions

■ Definition

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \left(\frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon} \right)$$

$$\frac{\partial f(x, y)}{\partial y} = \lim_{\varepsilon \rightarrow 0} \left(\frac{f(x, y + \varepsilon) - f(x, y)}{\varepsilon} \right)$$

■ Approximation

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x_{n+1}, y_m) - f(x_n, y_m)}{\Delta x}$$

$$\frac{\partial f(x, y)}{\partial y} \approx \frac{f(x_n, y_{m+1}) - f(x_n, y_m)}{\Delta x}$$

■ Convolution kernels

$$f_x = [1 \quad -1]$$

$$f_y = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

Discrete Edge Operators


How can we differentiate a *digital* image?

Finite difference approximations:

Gradients:

$$\frac{\partial I}{\partial x} \approx \frac{1}{2\varepsilon} ((I_{i+1,j+1} - I_{i,j+1}) + (I_{i+1,j} - I_{i,j}))$$

$$\frac{\partial I}{\partial y} \approx \frac{1}{2\varepsilon} ((I_{i+1,j+1} - I_{i+1,j}) + (I_{i,j+1} - I_{i,j}))$$



Convolution (cross-correlation) masks :

$$\frac{\partial I}{\partial x} \approx \frac{1}{2\varepsilon} \begin{array}{|c|c|} \hline -1 & 1 \\ \hline -1 & 1 \\ \hline \end{array}$$

$$\frac{\partial I}{\partial y} \approx \frac{1}{2\varepsilon} \begin{array}{|c|c|} \hline 1 & 1 \\ \hline -1 & -1 \\ \hline \end{array}$$

Discrete Edge Operators

2nd order partial derivatives:

$$\frac{\partial^2 I}{\partial x^2} \approx \frac{1}{\varepsilon^2} (I_{i-1,j} - 2I_{i,j} + I_{i+1,j})$$

$$\frac{\partial^2 I}{\partial y^2} \approx \frac{1}{\varepsilon^2} (I_{i,j-1} - 2I_{i,j} + I_{i,j+1})$$

$I_{i-1,j+1}$	$I_{i,j+1}$	$I_{i+1,j+1}$
$I_{i-1,j}$	$I_{i,j}$	$I_{i+1,j}$
$I_{i-1,j-1}$	$I_{i,j-1}$	$I_{i+1,j-1}$

Laplacian :


$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

Convolution (cross-correlation) masks :

$$\nabla^2 I \approx \frac{1}{\varepsilon^2} \begin{array}{|c|c|c|} \hline 0 & 1 & 0 \\ \hline 1 & -4 & 1 \\ \hline 0 & 1 & 0 \\ \hline \end{array}$$

or $\frac{1}{6\varepsilon^2}$

$$\begin{array}{|c|c|c|} \hline 1 & 4 & 1 \\ \hline 4 & -20 & 4 \\ \hline 1 & 4 & 1 \\ \hline \end{array}$$



The Sobel Operator

- Better approximations of the gradients exist
 - The *Sobel*/operators below are very commonly used


$$\frac{1}{8} \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array} s_x$$

$$\frac{1}{8} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array} s_y$$

- The standard defn. of the Sobel operator omits the 1/8 term
 - doesn't make a difference for edge detection
 - the 1/8 term **is** needed to get the right gradient value


Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal




- Where is the edge?


Solution: smooth first




h



$h \star f$



$\frac{\partial}{\partial x}(h \star f)$




Where is the edge? Look for peaks in $\frac{\partial}{\partial x}(h \star f)$ ³⁹


Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h * f) = (\frac{\partial}{\partial x}h) * f \quad \text{This saves us one operation:}$$


f



$\frac{\partial}{\partial x}h$



$(\frac{\partial}{\partial x}h) * f$



Laplacian of Gaussian

$$\frac{\partial^2}{\partial x^2}(h * f) = \left(\frac{\partial^2}{\partial x^2} h \right) * f$$

Laplacian of Gaussian




Sigma = 50

Consider $\frac{\partial^2}{\partial x^2}(h * f)$

f

$\frac{\partial^2}{\partial x^2} h$



$(\frac{\partial^2}{\partial x^2} h) * f$




Where is the edge?

Zero-crossings of bottom graph

2D edge detection filters




$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$




$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian




- ∇^2 is the **Laplacian** operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$



Edge Detection

- Prewitt and Sobel edge detectors
 - Compute derivatives
 - In x and y directions
 - Find gradient magnitude
 - Threshold gradient magnitude
- Difference between Prewitt and Sobel is the derivative filters



Prewitt Edge Detector

Prewitt's
edges in x
direction

$$\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

 I_x

Prewitt's
edges in y
direction

$$\begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

 I_y

Sobel Edge Detector




Sobel's
edges in x
direction

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$



I_x





Sobel's
edges in y
direction

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$





I_y







Prewitt Edge Detector






Sobel Edge Detector






Sobel Edge Detector



Edge Finding: Matlab Demo



```
im = imread('football.jpg');
image(im);
figure(2);
bw = double(rgb2gray(im));

[dx,dy] = gradient(bw);
gradmag = sqrt(dx.^2 + dy.^2);
image(gradmag);
```







Image Smoothing With Gaussian




```
figure(3);
sigma = 3;
width = 3 * sigma;
support = -width : width;
gauss2D = exp( - (support / sigma).^2 / 2);
gauss2D = gauss2D / sum(gauss2D);
smooth = conv2(bw, gauss2D, 'same');
image(smooth);
colormap(gray(255));
tic ; smooth = conv2(bw,gauss3D, 'same'); toc
```


Edge Detection With Smoothed Images



```
figure(4);
[dx,dy] = gradient(smooth);
gradmag = sqrt(dx.^2 + dy.^2);
gmax = max(max(gradmag));
imshow(gradmag);
colormap(gray(gmax));
```





Displaying the Edge Normal



```
figure(5);
hold on;
image(smooth);
colormap(gray(255));
[m,n] = size(gradmag);

edges = (gradmag > 0.3 * gmax);
inds = find(edges);
[posx, posy] = meshgrid(1:n, 1:m); posx2= posx(inds); posy2= posy(inds);
gm2= gradmag(inds);
sintheta = dx(inds) ./ gm2;
costheta = - dy(inds) ./ gm2;
quiver(posx2, posy2, gm2 .* sintheta / 10, -gm2 .* costheta / 10, 0);
hold off;
```






Sobel Operator

$$S_1 = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

$$S_2 = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$


$$\text{Edge Magnitude} = \sqrt{S_1^2 + S_2^2}$$

$$\text{Edge Direction} = \tan^{-1} \left(\frac{S_1}{S_2} \right)$$



Sobel Edge Detector


```
figure(6)  
edge(bw, 'sobel')
```




Edge detection

- Matlab demo


```
I = imread('circuit.tif');  
imshow(I);  
BW1 = edge(I,'prewitt');  
BW2 = edge(I,'canny');  
Figure;  
imshow(BW1);  
Figure;  
imshow(BW2);
```



Original image



Prewitt filter



Canny filter

Edge detection

- Matlab demo

```
I = imread('coins.png');  
imshow(I);  
BW1 = edge(I,'sobel');  
BW2 = edge(I,'canny');
```

Figure;


```
imshow(BW1);
```

Figure;


```
imshow(BW2);
```




Original image



Sobel filter



Canny filter




Features in Matlab

`edge(im,'prewitt')` - (almost) linear

`edge(im,'sobel')` - (almost) linear

`edge(im,'canny')` - not local, no closed form



Sobel Operator


$$S_1 = \begin{array}{|c|c|c|} \hline -1 & -2 & -1 \\ \hline 0 & 0 & 0 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

$$S_2 = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$


$$\text{Edge Magnitude} = \sqrt{S_1^2 + S_2^2}$$

$$\text{Edge Direction} = \tan^{-1} \left(\frac{S_1}{S_2} \right)$$

Sobel filter




`edge(im,'sobel')`



Today's Goals



- Features Overview
- Canny Edge Detector
- Harris Corner Detector
- Templates and Image Pyramid
- SIFT Features



Canny Edge Detector


- J. Canny, "A computational approach to edge detection, " IEEE Transactions on Pattern Analysis and Machine Intelligence, vol. 8, pp. 679--698, 1986
- Source code:
 - ftp://figment.csee.usf.edu/pub/Edge_Comparison/source_code/canny.src

Canny Edge Detector




`edge(im,'canny')`


Comparison



Sobel




Canny




Optimal Edge Detection: Canny

- Assume:
 - Linear filtering
 - Additive iid Gaussian noise
- Edge detector should have:
 - Good Detection. Filter responds to edge, not noise.
 - Good Localization: detected edge near true edge.
 - Single Response: one per edge.




Optimal Edge Detection: Canny (continued)

- Optimal Detector is approximately Derivative of Gaussian.
- Detection/Localization trade-off
 - More smoothing improves detection
 - And hurts localization.
- This is what you might guess from (detect change) + (remove noise)



Canny Edge Detector

- Criterion 1: Good Detection: The optimal detector must minimize the probability of false positives as well as false negatives.
- Criterion 2: Good Localization: The edges detected must be as close as possible to the true edges.
- Single Response Constraint: The detector must return one point only for each edge point.



Canny Edge Detector Steps

1. Smooth image with Gaussian filter
2. Compute derivative of filtered image
3. Find magnitude and orientation of gradient
4. Apply “Non-maximum Suppression”
5. Apply “Hysteresis Threshold”

Canny Edge Detector

First Two Steps

1. Filter out noise
 - Use a 2D Gaussian Filter. $J = I \otimes G$
2. Take a derivative
 - Compute the magnitude of the gradient:

$\nabla J = (J_x, J_y) = \left(\frac{\partial J}{\partial x}, \frac{\partial J}{\partial y} \right)$ is the Gradient

$$\|\nabla J\| = \sqrt{J_x^2 + J_y^2}$$

Canny Edge Detector

First Two Steps

- Smoothing

$$S = I * g(x, y) = g(x, y) * I$$

$$g(x, y) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

- Derivative

$$\nabla S = \nabla(g * I) = (\nabla g) * I$$





$$\nabla S = \begin{bmatrix} g_x \\ g_y \end{bmatrix} * I = \begin{bmatrix} g_x * I \\ g_y * I \end{bmatrix}$$


Homework

$$\nabla g = \begin{bmatrix} \frac{\partial g}{\partial x} \\ \frac{\partial g}{\partial y} \end{bmatrix} = \begin{bmatrix} g_x \\ g_y \end{bmatrix}$$

Canny Edge Detector



Derivative of Gaussian







Smoothing and Differentiation

- Need two derivatives, in x and y direction.
- We can use a derivative of Gaussian filter
 - because differentiation is convolution, and convolution is associative



Canny Edge Detector

First Two Steps



Canny Edge Detector

Third Step

- Gradient magnitude and gradient direction

(S_x, S_y) Gradient Vector

$$\text{magnitude} = \sqrt{(S_x^2 + S_y^2)}$$


$$\text{direction} = \theta = \tan^{-1} \frac{S_y}{S_x}$$



image



gradient magnitude





Finding the Peak

- 1) The gradient magnitude is large along thick trail; how do we identify the significant points?
- 2) How do we link the relevant points up into curves?


Canny Edge Detector

Fourth Step



- Non maximum suppression



We wish to mark points along the curve where the **magnitude is biggest**. We can do this by looking for a maximum along a slice normal to the curve (non-maximum suppression). These points should form a curve. There are then two algorithmic issues: at which point is the maximum, and where is the next one?



Non-Maximum Supression



Non-maximum suppression:

Select the single maximum point across the width of an edge.



p



q



r



Gradient



(Forsyth & Ponce)

Non-maximum suppression


At q, we have a maximum if the value is larger than those at both p and at r. Interpolate to get these values.




Predicting
the next
edge point

Assume the
marked point is an
edge point. Then
we construct the
tangent to the edge
curve (which is
normal to the
gradient at that
point) and use this
to predict the next
points (here either
 r or s).

Gradient




(Forsyth & Ponce)



Canny Edge Detector

Non-Maximum Suppression

- Suppress the pixels in $|\nabla S|$ which are not local maximum



$$M(x, y) = \begin{cases} |\nabla S|(x, y) & \text{if } |\Delta S|(x, y) > |\Delta S|(x', y') \\ & \quad \& |\Delta S|(x, y) > |\Delta S|(x'', y'') \\ 0 & \text{otherwise} \end{cases}$$

x' and x'' are the neighbors of x along normal direction to an edge

Canny Edge Detector

Quantization of Normal Directions

$$\tan \theta = \frac{S_y}{S_x}$$


Quantizations :

0: $-0.4142 < \tan \theta \leq 0.4142$

1: $0.4142 < \tan \theta < 2.4142$

2: $|\tan \theta| \geq 2.4142$

3: $-2.4142 < \tan \theta \leq -0.4142$



Canny Edge Detector

Non-Maximum Suppression

$$|\Delta S| = \sqrt{S_x^2 + S_y^2}$$




M



For visualization


$$M \geq \text{Threshold} = 25$$






Hysteresis(滯变)

- Check that maximum value of gradient value is sufficiently large
 - drop-outs? use **hysteresis**
 - use a high threshold to start edge curves and a low threshold to continue them.






Edge Hysteresis

- **Hysteresis:** A lag or momentum factor
- Idea: Maintain two thresholds k_{high} and k_{low}
 - Use k_{high} to find strong edges to start edge chain
 - Use k_{low} to find weak edges which continue edge chain
- Typical ratio of thresholds is roughly

$$k_{\text{high}} / k_{\text{low}} = 2$$



Canny Edge Detector


Hysteresis Thresholding

- If the gradient at a pixel is
 - above “**High**”, declare it an ‘**edge pixel**’
 - below “**Low**”, declare it a “**non-edge-pixel**”
 - **between** “low” and “high”
 - Consider its neighbors iteratively then declare it an “edge pixel” if it is **connected** to an ‘edge pixel’ **directly** or via pixels **between** “low” and “high”.


Canny Edge Detector

Hysteresis Thresholding


- Connectedness



4 connected





8 connected



6 connected

Canny Edge Detector Hysteresis Thresholding






Canny Edge Detector

Hysteresis Thresholding

- Scan the image from left to right, top-bottom.
 - The gradient magnitude at a pixel is above a high threshold declare that as an edge point
 - Then recursively consider the *neighbors* of this pixel.
 - If the gradient magnitude is above the low threshold declare that as an edge pixel.

Canny Edge Detector Hysteresis Thresholding



Canny Edge Detector Hysteresis Thresholding

M



regular


$M \geq 25$



Hysteresis

High = 35

Low = 15



Summary: Canny Edge Detector


Steps:

1. Apply derivative of Gaussian
2. Non-maximum suppression
 - Thin multi-pixel wide “ridges” down to single pixel width
3. Linking and thresholding
 - Low, high edge-strength thresholds
 - Accept all edges over low threshold that are connected to edge over high threshold

Summary: Canny Edge Operator

- Smooth image I with 2D Gaussian: $G * I$
- Find local edge normal directions for each pixel $\bar{\mathbf{n}} = \frac{\nabla(G * I)}{|\nabla(G * I)|}$
- Compute edge magnitudes $|\nabla(G * I)|$
- Find the location of the edges by finding zero-crossings along the edge normal directions (**non-maximum suppression**) $\frac{\partial^2(G * I)}{\partial \bar{\mathbf{n}}^2} = 0$
- Threshold edges in the image with **hysteresis** to eliminate spurious responses

Read Canny's original paper for further details



Why is Canny so Dominant

- Still widely used after 20 years.
 1. Theory is nice (but end result same).
 2. Details good (magnitude of gradient).
 3. Hysteresis an important heuristic.
 4. Code was distributed.
 5. Perhaps this is about all you can do with linear filtering.



Demo of Edge Detection

Canny Edge Detection (Example)


Original image




Strong edges only



Weak edges




Canny Edge Detection (Example)




Using Matlab with default thresholds

Bridge Example




`edge(im,'canny')`



The Canny Edge Detector



original image (Lena)



The Canny Edge Detector




magnitude of the gradient

The Canny Edge Detector




After non-maximum suppression and thresholding with hysteresis

Canny Edge Operator




- The choice of σ depends on desired behavior
 - large σ detects large scale edges
 - small σ detects fine features






**fine scale
high
threshold**




coarse
scale,
high
threshold





A binary image showing a brain segmentation. The image consists of black outlines on a white background, representing different tissue types or regions. A legend is located in the bottom right corner, consisting of four horizontal lines followed by the text "coarse scale low threshold".

coarse
scale
low
threshold




Corner Effects






Today's Goals (Break)


- Features Overview
- Canny Edge Detector
- Harris Corner Detector
- Templates and Image Pyramid
- SIFT Features



Corners(Start)


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners(End)


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners


- Why are they important?






Corners(End)


- Why are they important?






Corners

- Why are they important?




Corners contain more edges than lines.

- A point on a line is hard to match.




Which one is the correct correspondence?



Corners contain more edges than lines.

- A corner is easier




Finding Corners

Edge detectors perform poorly at corners.



Corners provide repeatable points for matching,
so are worth detecting.

Idea:

- Right at a corner, gradient is ill defined.
- Near a corner, gradient has two or more different values.



Edge Detectors Tend to Fail at Corners



$$I_x = ?$$

$$I_y = ?$$

Formula for Finding Corners

Look at the second-moment matrix:

Sum over a small region,
the hypothetical corner



$$C = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix}$$

Gradient with respect to x,
times gradient with respect to y

Matrix is symmetric


WHY THIS? 125

Simple Case




First, consider case where:

$$C = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$




This means dominant gradient directions align with x or y axis



If either λ is close to 0, then this is **not** a corner, so look for locations where both are large.




General Case



Rotate



Shear




General Case

It can be shown that since C is rotationally symmetric:

$$C = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$


So every case is like a rotated version of the standard one on last slide.




So, to detect corners

- Filter image.
- Compute magnitude of the gradient everywhere.
- Construct C in a window around the target pixel.
- Use Linear Algebra to find λ_1 and λ_2 .
- If they are both big, we have a corner.

Gradient Orientation




Corner Detection



Corners are detected where the product of the ellipse axis lengths reaches a local maximum.



Harris Corners




- Originally developed as features for motion tracking
- Greatly reduces amount of computation compared to tracking every pixel
- Translation and rotation invariant (but not scale invariant)


Harris Corner: Matlab code

```
% Harris Corner detector - by Kashif Shahzad
sigma=2; thresh=0.1; sze=11; disp=0;eps=0.0;
dy = [-1 0 1; -1 0 1; -1 0 1]; % Derivative masks
dx = dy'; %dx is the transpose matrix of dy
% Ix and Iy are the horizontal and vertical edges of image
I = imread('rice.png');
imshow(I);
title('\bf Original image');%use bold font for the title
bw=double(I);%convert uint8 to double
Ix = conv2(bw, dx, 'same'); % Calculating the gradient
Iy = conv2(bw, dy, 'same'); %return a matrix the same
g = fspecial('gaussian',max(1,fix(6*sigma)), sigma); %
Ix2 = conv2(Ix.^2, g, 'same'); %Smoothed squared image
Iy2 = conv2(Iy.^2, g, 'same');
Ixy = conv2(Ix.*Iy, g, 'same');
cornerness = (Ix2.*Iy2 - Ixy.^2)./(Ix2 + Iy2 + eps); %
mx = ordfilt2(cornerness,sze^2,ones(sze)); %
cornerness = (cornerness==mx)&(cornerness>thresh); %
[rws,cols] = find(cornerness); %
figure;imshow(bw,[0 255]);
hold on;
p=[cols rws];
plot(p(:,1),p(:,2),'or'); % display corners as red cir
title('\bf Harris Corners');
```







Example ($\sigma=0.1$)




Example ($\sigma=0.01$)





Example ($\sigma=0.001$)




Reading: Matching with Invariant Features

(www.cs.washington.edu, computer vision course)


Harris corner detector

- C.Harris, M.Stephens. “A Combined Corner and Edge Detector”. 1988




The Basic Idea


- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity




Harris Detector: Basic Idea



“flat” region:
no change in
all directions



“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions

Harris Detector: Mathematics

Change of intensity for the shift $[u, v]$:


$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window
function

Shifted
intensity


Intensity

Window function $w(x, y) =$



1 in window, 0 outside

or



Gaussian

Harris Detector: Mathematics

For small shifts $[u, v]$ we have a *bilinear* approximation:

$$E(u, v) \cong [u, v] \ M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Harris Detector: Mathematics

Intensity change in shifting window: eigenvalue analysis


$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

λ_1, λ_2 – eigenvalues of M

direction of the
fastest change

Ellipse $E(u, v) = \text{const}$

direction of the
slowest change





Why? Optional assignment


Harris Detector: Mathematics

Classification of image points using eigenvalues of M :

λ_2

λ_1 and λ_2 are small;
 E is almost constant
in all directions






Harris Detector: Mathematics

Measure of corner response:

$$R = \det M - k (\operatorname{trace} M)^2$$


$$\det M = \lambda_1 \lambda_2$$
$$\operatorname{trace} M = \lambda_1 + \lambda_2$$


(k – empirical constant, $k = 0.04\text{-}0.06$)



Harris Detector: Mathematics

- R depends only on eigenvalues of \mathbf{M}
- R is large for a corner
- R is negative with large magnitude for an edge
- $|R|$ is small for a flat region





Harris Detector


- The Algorithm:
 - Find points with large corner response function R ($R >$ threshold)
 - Take the points of local maxima of R

Harris Detector: Workflow




Harris Detector: Workflow

Compute corner response R




Harris Detector: Workflow

Find points with large corner response: $R > \text{threshold}$




Harris Detector: Workflow

Take only the points of local maxima of R



Harris Detector: Workflow





Harris Detector: Summary


- Average intensity change in direction $[u, v]$ can be expressed as a bilinear form:

$$E(u, v) \cong [u, v] \ M \begin{bmatrix} u \\ v \end{bmatrix}$$

- Describe a point in terms of eigenvalues of M :
measure of corner response
- A good (corner) point should have a *large intensity change in all directions*, i.e. R should be large positive

Harris Detector: Some Properties

- Rotation invariance





Ellipse rotates but its shape (i.e. eigenvalues)
remains the same

Corner response R is invariant to image rotation


Harris Detector: Some Properties

- Partial invariance to *affine intensity* change
 - ✓ Only derivatives are used => invariance to intensity shift: $I \rightarrow I + b$
 - ✓ Intensity scale: $I \rightarrow a I$



Harris Detector: Some Properties

But: non-invariant to *image scale!*



All points will be
classified as edges

Corner !



Models of Image Change


Geometry

- Rotation
- Similarity (rotation + uniform scale)
- Affine (scale dependent on direction)
valid for: orthographic camera, locally planar object

Photometry


- Affine intensity change ($I \rightarrow aI + b$)






Scale Invariant Detection


- Consider regions (e.g. circles) of different sizes around a point
- Regions of corresponding sizes will look the same in both images






Scale Invariant Detection

- The problem: how do we choose corresponding circles *independently* in each image?






Today's Goals

- Features Overview
- Canny Edge Detector
- Harris Corner Detector
- Templates and Image Pyramid
- SIFT Features

Problem: Features for Recognition

Want to find



... in here



Correlation(相关)



template

How do we locate the template in the image?


Minimize

$$\begin{aligned} E(i, j) &= \sum \sum [f(m, n) - t(m - i, n - j)]^2 \\ &= \sum_m^m \sum_n^n [f^2(m, n) + t^2(m - i, n - j) - 2f(m, n)t(m - i, n - j)] \end{aligned}$$

Maximize

$$R_{tf}(i, j) = \sum_m \sum_n t(m - i, n - j)f(m, n)$$

Cross-correlation



Correlation (相关)

- Cauchy inequality (柯西不等式)

$$a^2 + b^2 + c^2 \geq ab + bc + ca$$

$$R\{(a,b,c), (a,b,c)\} > R\{(a,b,c), (b,c,a)\}$$

$$4a^2 + 4b^2 + 4c^2 \geq a^2 + b^2 + c^2$$

$$R\{(a,b,c), (4a,4b,4c)\} > R\{(a,b,c), (a,b,c)\} ?$$



Cross-correlation (互相关)

$$R_{tf}(i, j) = \sum_m \sum_n t(m-i, n-j) f(m, n) \quad R_{tf} = t \otimes f$$

Note: $t \otimes f \neq f \otimes t$


$$R_{ff} = f \otimes f \quad \text{Auto-correlation}$$

Problem:



$$R_{tf}(C) > R_{tf}(B) > R_{tf}(A)$$

We need $R_{tf}(A)$ to be the maximum!¹⁶⁴



Correlation

- Cauchy inequality (柯西不等式)


$$\text{Corr}(A, B) = \text{dot}(A, B) / \sqrt{|A| |B|}$$

$$\text{Corr}\{(a,b,c), (4a,4b,4c)\} = \text{Corr}\{(a,b,c), (a,b,c)\} = 1.0$$

Normalized Correlation

Account for energy differences


$$N_{tf}(i, j) = \frac{\sum_m \sum_n t(m-i, n-j) f(m, n)}{\left[\sum_m \sum_n t^2(m-i, n-i) \right]^{1/2} \left[\sum_m \sum_n f^2(m, n) \right]^{1/2}}$$



Normalized Correlation

```
onion = imread('onion.png');
peppers = imread('peppers.png');
imshow(onion);
figure, imshow(peppers);
rect_onion = [111 33 65 58];
rect_peppers = [163 47 143 151];
sub_onion = imcrop(onion,rect_onion);
sub_peppers = imcrop(peppers,rect_peppers);
c = normxcorr2(sub_onion(:,:,1),sub_peppers(:,:,1));
[max_c, imax] = max(abs(c(:)));
[ypeak, xpeak] = ind2sub(size(c),imax(1));
corr_offset = [(xpeak-size(sub_onion,2)); (ypeak-size(sub_onion,1))];
rect_offset = [(rect_peppers(1)-rect_onion(1)); (rect_peppers(2)-rect_onion(2))];
offset = corr_offset + rect_offset;
xoffset = offset(1);
yoffset = offset(2);
xbegin = round(xoffset+1);
xend = round(xoffset+ size(onion,2));
ybegin = round(yoffset+1);
yend = round(yoffset+size(onion,1));
extracted_onion = peppers(ybegin:yend,xbegin:xend,:);
recovered_onion = uint8(zeros(size(peppers)));
recovered_onion(ybegin:yend,xbegin:xend,:) = onion;
[m,n,p] = size(peppers);
mask = ones(m,n);
i = find(recovered_onion(:,:1)==0);
mask(i) = .2;
figure, imshow(peppers(:,:,:));
hold on;
h = imshow(recovered_onion);
set(h,'AlphaData',mask);
```







Templates


- Find an object in an image!
- Want Invariance!
 - Scaling
 - Rotation
 - Illumination
 - Deformation

Template Convolution



Template Convolution







Convolution with Templates

- Invariances:
 - Scaling No
 - Rotation No
 - Illumination No
 - Deformation Maybe
- Provides
 - Good localization No


Scale Invariance: Image Pyramid







Templates with Image Pyramid

- Invariances:
 - Scaling Yes
 - Rotation No
 - Illumination No
 - Deformation Maybe
- Provides
 - Good localization No




Templates





Optional Assignment— Feature detector


- Point feature detector
- Line feature detector
- Conic feature detector
- Invariance under different cases
- Feature matching/Correspondence
-




Today's Goals

- Canny Edge Detector
- Harris Corner Detector
- Hough Transform
- Templates and Image Pyramid
- SIFT Features

This is the end of features....




Appendix



SIFT

- Invariances:
 - Scaling Yes
 - Rotation Yes
 - Illumination Yes
 - Deformation Maybe
- Provides
 - Good localization Yes


Invariance to ...



Scaling and rotation

Viewpoint

Illumination




SIFT Reference

Distinctive image features from scale-invariant keypoints. David G. Lowe, International Journal of Computer Vision, 60, 2 (2004), pp. 91-110.


SIFT = Scale Invariant Feature Transform

Invariant Local Features

- Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters



Advantages of invariant local features




- **Locality:** features are local, so robust to occlusion and clutter (no prior segmentation)
- **Distinctiveness:** individual features can be matched to a large database of objects
- **Quantity:** many features can be generated for even small objects
- **Efficiency:** close to real-time performance
- **Extensibility:** can easily be extended to wide range of differing feature types, with each adding robustness

SIFT On-A-Slide

1. **Enforce invariance to scale:** Compute Gaussian difference max, for many different scales; non-maximum suppression, find local maxima: keypoint candidates
2. **Localizable corner:** For each maximum fit quadratic function. Compute center with sub-pixel accuracy by setting first derivative to zero.
3. **Eliminate edges:** Compute ratio of eigenvalues, drop keypoints for which this ratio is larger than a threshold.
4. **Enforce invariance to orientation:** Compute orientation, to achieve scale invariance, by finding the strongest second derivative direction in the smoothed image (possibly multiple orientations). Rotate patch so that orientation points up.
5. **Compute feature signature:** Compute a "gradient histogram" of the local image region in a 4x4 pixel region. Do this for 4x4 regions of that size. Orient so that largest gradient points up (possibly multiple solutions). Result: feature vector with 128 values (15 fields, 8 gradients).
6. **Enforce invariance to illumination change and camera saturation:** Normalize to unit length to increase invariance to illumination. Then threshold all gradients, to become invariant to camera saturation.

SIFT On-A-Slide

1. **Enforce invariance to scale:** Compute Gaussian difference max, for many different scales; non-maximum suppression, find local maxima: keypoint candidates
2. **Localizable corner:** For each maximum fit quadratic function. Compute center with sub-pixel accuracy by setting first derivative to zero.
3. **Eliminate edges:** Compute ratio of eigenvalues, drop keypoints for which this ratio is larger than a threshold.
4. **Enforce invariance to orientation:** Compute orientation, to achieve scale invariance, by finding the strongest second derivative direction in the smoothed image (possibly multiple orientations). Rotate patch so that orientation points up.
5. **Compute feature signature:** Compute a "gradient histogram" of the local image region in a 4x4 pixel region. Do this for 4x4 regions of that size. Orient so that largest gradient points up (possibly multiple solutions). Result: feature vector with 128 values (15 fields, 8 gradients).
6. **Enforce invariance to illumination change and camera saturation:** Normalize to unit length to increase invariance to illumination. Then threshold all gradients, to become invariant to camera saturation.




Finding “Keypoints” (Corners)

Idea: Find Corners, but scale invariance


Approach:

- Run linear filter (diff of Gaussians)
- At different resolutions of image pyramid


Difference of Gaussians




Minus



Equals






Difference of Gaussians

```
surf(fspecial('gaussian',40,4))
```

```
surf(fspecial('gaussian',40,8))
```

```
surf(fspecial('gaussian',40,8) - fspecial('gaussian',40,4))
```




Find Corners with DiffOfGauss


```
im = imread('bridge.jpg');
bw = double(im(:,:,1)) / 256;

for i = 1 : 10
    gaussD = fspecial('gaussian',40,2*i) -
        fspecial('gaussian',40,i);
    res = abs(conv2(bw, gaussD, 'same'));
    res = res / max(max(res));
    imshow(res) ; title(['\bf i = ' num2str(i)]); drawnow
end
```


Gaussian Kernel Size i=1



Gaussian Kernel Size i=2



Gaussian Kernel Size i=3




Gaussian Kernel Size i=4




Gaussian Kernel Size i=5




Gaussian Kernel Size i=6




Gaussian Kernel Size i=7




Gaussian Kernel Size i=8




Gaussian Kernel Size i=9




Gaussian Kernel Size $i=10$





Key point localization

- Detect maxima and minima of difference-of-Gaussian in scale space



Example of keypoint detection



- (a)** 233x189 image
- (b)** 832 DOG extrema
- (c)** 729 above threshold

SIFT On-A-Slide

1. **Enforce invariance to scale:** Compute Gaussian difference max, for many different scales; non-maximum suppression, find local maxima: keypoint candidates
2. **Localizable corner:** For each maximum fit quadratic function. Compute center with sub-pixel accuracy by setting first derivative to zero.
3. **Eliminate edges:** Compute ratio of eigenvalues, drop keypoints for which this ratio is larger than a threshold.
4. **Enforce invariance to orientation:** Compute orientation, to achieve scale invariance, by finding the strongest second derivative direction in the smoothed image (possibly multiple orientations). Rotate patch so that orientation points up.
5. **Compute feature signature:** Compute a "gradient histogram" of the local image region in a 4x4 pixel region. Do this for 4x4 regions of that size. Orient so that largest gradient points up (possibly multiple solutions). Result: feature vector with 128 values (15 fields, 8 gradients).
6. **Enforce invariance to illumination change and camera saturation:** Normalize to unit length to increase invariance to illumination. Then threshold all gradients, to become invariant to camera saturation.

Example of keypoint detection

Threshold on value at DOG peak and on ratio of principle curvatures
(Harris approach)




(c) 729 left after peak value threshold (from 832)

(d) 536 left after testing ratio of principle curvatures



SIFT On-A-Slide

1. **Enforce invariance to scale:** Compute Gaussian difference max, for many different scales; non-maximum suppression, find local maxima: keypoint candidates
2. **Localizable corner:** For each maximum fit quadratic function. Compute center with sub-pixel accuracy by setting first derivative to zero.
3. **Eliminate edges:** Compute ratio of eigenvalues, drop keypoints for which this ratio is larger than a threshold.
4. **Enforce invariance to orientation:** Compute orientation, to achieve scale invariance, by finding the strongest second derivative direction in the smoothed image (possibly multiple orientations). Rotate patch so that orientation points up.
5. **Compute feature signature:** Compute a "gradient histogram" of the local image region in a 4x4 pixel region. Do this for 4x4 regions of that size. Orient so that largest gradient points up (possibly multiple solutions). Result: feature vector with 128 values (15 fields, 8 gradients).
6. **Enforce invariance to illumination change and camera saturation:** Normalize to unit length to increase invariance to illumination. Then threshold all gradients, to become invariant to camera saturation.



Select canonical orientation

- Create histogram of local gradient directions computed at selected scale
- Assign canonical orientation at peak of smoothed histogram
- Each key specifies stable 2D coordinates (x , y , scale, orientation)



SIFT On-A-Slide

1. **Enforce invariance to scale:** Compute Gaussian difference max, for many different scales; non-maximum suppression, find local maxima: keypoint candidates
2. **Localizable corner:** For each maximum fit quadratic function. Compute center with sub-pixel accuracy by setting first derivative to zero.
3. **Eliminate edges:** Compute ratio of eigenvalues, drop keypoints for which this ratio is larger than a threshold.
4. **Enforce invariance to orientation:** Compute orientation, to achieve scale invariance, by finding the strongest second derivative direction in the smoothed image (possibly multiple orientations). Rotate patch so that orientation points up.
5. **Compute feature signature:** Compute a "gradient histogram" of the local image region in a 4x4 pixel region. Do this for 4x4 regions of that size. Orient so that largest gradient points up (possibly multiple solutions). Result: feature vector with 128 values (15 fields, 8 gradients).
6. **Enforce invariance to illumination change and camera saturation:** Normalize to unit length to increase invariance to illumination. Then threshold all gradients, to become invariant to camera saturation.

SIFT vector formation

- Thresholded image gradients are sampled over 16x16 array of locations in scale space
- Create array of orientation histograms
- 8 orientations x 4x4 histogram array = 128 dimensions






Image gradients



Keypoint descriptor

Nearest-neighbor matching to feature database

- Hypotheses are generated by **approximate nearest neighbor** matching of each feature to vectors in the database
 - SIFT use best-bin-first (Beis & Lowe, 97) modification to k-d tree algorithm
 - Use heap data structure to identify bins in order by their distance from query point
- **Result:** Can give speedup by factor of 1000 while finding nearest neighbor (of interest) 95% of the time



3D Object Recognition



- Extract outlines with background subtraction



3D Object Recognition



- Only 3 keys are needed for recognition, so extra keys provide robustness
- Affine model is no longer as accurate




Recognition under occlusion






Test of illumination invariance

- Same image under differing illumination





273 keys verified in final match

Examples of view interpolation



Location recognition





SIFT

- Invariances:
 - Scaling Yes
 - Rotation Yes
 - Illumination Yes
 - Deformation Maybe
- Provides
 - Good localization Yes